

NEURO-CIRCT: Neuro-Symbolic CEGIS with SSA Provenance Slicing for Autonomous Hardware Compiler Repair

Kudchadkar

CHIA Hackathon 2026 • Track §5.5: Autonomous CIRCT Issue Solving • A³ Workshop at MICRO 2026
Artifact Repository: <https://github.com/kudchadkar/circt-context-precision-gate>

Abstract—Automated program repair (APR) for extensible intermediate representation (IR) compilers presents unique challenges at the intersection of program analysis, type theory, and formal verification. In multi-dialect compiler infrastructures such as LLVM/CIRCT, defect triggers frequently manifest as monolithic MLIR modules comprising thousands of operations, while regression validation over full suites incurs prohibitive turnaround latencies. Furthermore, generative code models exhibit severe failure modes on compiler codebases—frequently suppressing diagnostic aborts by deleting invariant assertions or introducing ill-typed early exits, rather than repairing the underlying dialect lowering. This paper presents **NEURO-CIRCT**, a neuro-symbolic framework for autonomous compiler repair that couples backward Static Single Assignment (SSA) def-use slicing, declarative TableGen/ODS reflection, multi-pass pipeline bisection, and Counterexample-Guided Inductive Synthesis (CEGIS) with bounded bit-vector equivalence validation (`circt-lec`). Evaluated across the 16-issue CHIA benchmark backlog (with external ground-truth triage labels from CIRCT maintainers) and focusing on three complex defect archetypes (including post-training-cutoff defect #10104), NEURO-CIRCT achieves **100% resolution within a 3-turn budget** (66.7% first-turn Pass@1) with zero assertion deletions. Integrating incremental C++ rebuilds (24.8s) and isolated dialect lit test execution (3.2s), NEURO-CIRCT achieves a **7.5× end-to-end turn speedup** (209.0s → 28.0s) and a **10.2× wall-clock acceleration to verified fix** (620.2s → 60.8s), while our cascaded SSA slicer reduces reproducers by up to 492.7× in 2.20s.

Index Terms—Hardware Compilers, CIRCT, MLIR, Automated Program Repair, Neuro-Symbolic Synthesis, Counterexample-Guided Inductive Synthesis, Program Slicing, Equivalence Validation, SMT.

I. INTRODUCTION

Modern Electronic Design Automation (EDA) flows increasingly rely on modular, extensible intermediate representations (IRs) to decouple high-level hardware description languages from backend physical synthesis. The LLVM CIRCT (Circuit Intermediate Representation for Compilers and Tools) framework, constructed atop MLIR [1], [2], provides an extensible multi-dialect infrastructure unifying digital circuit elaboration, hardware dialect lowering, formal verification, and SystemVerilog emission across dialects including FIRRTL, HW, Comb, and SMT.

While automated program repair (APR) powered by Large Language Models (LLMs) has demonstrated substantial progress on standard software engineering benchmarks such as SWE-bench [3], applying autonomous repair agents to production hardware compilers exposes three fundamental domain-specific challenges:

- **Context Overhead from Monolithic IR Modules:** Hardware compiler reproducers routinely exceed 5,000 nested MLIR operations with deeply interconnected value definitions and multi-dialect type hierarchies. Presenting unreduced modules directly to neural synthesis models dilutes causal attention across thousands of extraneous operations, degrading defect localization and triggering hallucinated dialect constructs.
- **High-Latency Regression Feedback Loops:** The CIRCT regression harness comprises over 2,000 integration tests executed via LLVM's `lit` framework. Running the complete suite incurs 3–5 minutes per compilation attempt, severely bottlenecking iterative repair loops and limiting hypothesis exploration within fixed compute budgets.
- **Unsound Defect Suppression Pathologies:** When confronted with assertion violations during pattern rewriting, unconstrained language models frequently generate unsound workarounds: deleting defensive `assert()` checks, injecting unconditional `return success()` exits without updating the IR graph, or performing unchecked pointer downcasts that violate LLVM RTTI conventions.

To systematically resolve these challenges in the context of the CHIA Autonomous CIRCT Issue Solving Track (§5.5), this paper presents **NEURO-CIRCT**, a neuro-symbolic automated program repair system specifically architected for multi-dialect hardware compilers. This paper makes the following contributions:

- **Causal SSA Provenance Slicing:** An inter-operational backward program slicing algorithm that traces dataflow dependencies from diagnostic abort sinks, reducing input module volume by up to 492.7×.
- **Multi-Pass Phase-Ordering Bisection:** A formal bisection procedure under monotonicity assumptions that isolates contract-violating predecessor passes in multi-stage compilation pipelines.
- **Declarative TableGen Semantic Reflection:** Runtime introspection of ODS specifications to constrain neural patch synthesis to valid

dialect traits, operand types, and verifier invariants.

- **Inductive Boundary Validation & Equivalence Checking:** A verification harness combining Counterexample-Guided Inductive Synthesis (CEGIS) [4] over boundary valuations with bounded transition miter verification via `circt-lec` and SMT solvers [5], [6].
- **Empirical Rigor & Generalization:** Evaluation on the 16-issue CHIA benchmark with external maintainer ground-truth triage labels, including post-training-cutoff defect #10104, demonstrating 100% verified fix resolution, maintainer-aligned compilable C++ patches, and a **7.5× end-to-end iteration speedup**.

II. PROBLEM FORMULATION & CORRECTNESS CRITERIA

Let M denote the universe of well-formed MLIR modules, and let $D = \{D_1, \dots, D_k\}$ denote a set of MLIR dialects defined by declarative TableGen specifications T_D . A compiler pass is a mapping $p: M \rightarrow M \cup \{\perp_{\text{crash}}, \perp_{\text{diag}}\}$, where $\perp_{\text{crash}}$ denotes an abnormal termination (assertion abort, crash, stack overflow) and $\perp_{\text{diag}}$ denotes a graceful diagnostic rejection or pattern match failure. A compilation pipeline $\Pi = p_n \circ \dots \circ p_1$ applies an ordered sequence of transformation passes to an initial module $m_0 \in M$.

A compiler defect is characterized by an input module m_{bug} such that $\Pi(m_{\text{bug}}) = \perp_{\text{crash}}$. The objective of automated compiler repair is to synthesize a source code patch Δ for the compiler codebase C yielding $C' = C \oplus \Delta$ such that four formal correctness invariants are satisfied:

1. **Dual-Mode Defect Resolution:** Fatal crashes must be eliminated: $\Pi_C(m_{\text{bug}}) \neq \perp_{\text{crash}}$. Specifically:
 - (a) *Legalization & Transformation (Listings 1–3):* If m_{bug} is a supported or legalizable input program, the pipeline produces a legalized module m_{out} satisfying dialect verifier $V_D(m_{\text{out}}) = \text{true}$ (e.g., zero-width aggregate constants in Listing 1, or scoped LTL expansion in Listing 3).
 - (b) *Graceful Diagnostic Rejection:* If m_{bug} represents an illegal or unsupported construct, the patched pipeline terminates cleanly via `rewriter.notifyMatchFailure` or `emitError` ($\perp_{\text{diag}}$).
2. **Semantic Preservation on Valid Programs:** For all programs m in the valid input language where the unpatched compiler completed without crashing ($\Pi_C(m) \notin \{\perp_{\text{crash}}, \perp_{\text{diag}}\}$), circuit semantics are preserved: $\llbracket \Pi_C(m) \rrbracket \equiv \llbracket \Pi_{C'}(m) \rrbracket$. Verified via full-suite regression differential testing.
3. **Invariant Monotonicity:** Defensive assertion checks are monotonic: no preexisting contract assertions $\phi \in \text{Inv}(C)$ may be deleted in C' .
4. **Syntactic & Compilation Integrity:** The patch Δ must compile cleanly with `-DLLVM_ENABLE_ASSERTIONS=ON` without errors and adhere to LLVM standards (2-space indent, canonical RTTI casting).

III. ARCHITECTURAL FRAMEWORK

NEURO-CIRCT decomposes the compiler repair workflow into eight modular stages (detailed in Algorithm 1):

A. Backward SSA Provenance Slicing

Compiler crashes are typically triggered by localized invalid dialect states rather than global module structure. Given a failing operation $O_{\text{fail}} \in M$, its backward def-use slice $S(O_{\text{fail}})$ is defined inductively over SSA values $v \in \text{Operands}(O_{\text{fail}})$:

$$S(O_{\text{fail}}) = \{O_{\text{fail}}\} \cup \bigcup_{v \in \text{Operands}(O_{\text{fail}})} S_{\text{val}}(v) \quad (1)$$

where $S_{\text{val}}(v) = \{\text{DefOp}(v)\} \cup S(\text{DefOp}(v))$. If v is a `BlockArgument`, its owner block is included as a boundary condition. For non-SSA constructs (symbols and memory effects), the slice is complemented by `circt-reduce`. Slicing alone strips 99.1% of operations in 0.38s ($5,420 \rightarrow 48$ ops), while cascaded reduction reaches 11 operations (**492.7× reduction**) in 2.20s.

B. Multi-Pass Pipeline Phase-Ordering Bisector

In multi-pass pipelines, downstream passes frequently abort because an upstream pass emitted an ill-formed intermediate representation that evaded verification. Formally, for pass sequence $\Pi = \langle p_1, \dots, p_n \rangle$ on IR state σ :

$$\exists k < m: \sigma_k \notin L(D) \wedge \text{Abort}(p_m(\sigma_{m-1})) \quad (2)$$

Assuming pass monotonicity (ill-formed states persist across subsequent lowering stages), `pipeline_bisector` executes binary prefix searches over Π with `--verify-each` enabled, isolating the earliest verifier-violating pass p_k .

C. Declarative TableGen / ODS Semantic Reflection

MLIR operations are formally declared in TableGen (.td) files via the Operation Definition Specification (ODS). The `tablegen_analyzer` parses ODS trait sets (e.g., `Pure`, `SameOperandsAndResultType`) and verifier constraints directly from the CIRCT source tree. By injecting these formal constraints into the synthesis context, the system eliminates hallucinated C++ methods and ensures rewrites adhere to dialect type contracts.

D. Fault Localization & Diagnostic Frame Slicing

Upon a compiler abort, `fault_localizer` parses the C++ traceback to extract the crash location (File, Line) (e.g., `LowerToHW.cpp:3572`). It constructs a bounded 30-line code window centered at the crash site, annotating the failure predicate with symbolic directional markers.

Algorithm 1: NEURO-CIRCT Autonomous Compiler Repair Loop

```

Input: Issue Spec  $I$ , Crash Repro  $R_{\text{raw}}$ , Dialects  $D$ , Compiler  $C$ 
Output: Verified Patch  $\Delta$ , Lit Test Guard  $T$ 
1:  $R_{\text{min}} \leftarrow \text{SSABackwardSlice}(R_{\text{raw}}) \circ \text{circt-reduce}(R_{\text{raw}})$ 
2:  $\Pi_{\text{root}} \leftarrow \text{BisectPipeline}(R_{\text{min}}, \text{passes})$ 
3:  $O_{\text{traits}} \leftarrow \text{TableGenODSReflect}(D, R_{\text{min}})$ 
4:  $L_{\text{crash}}, C_{\text{slice}} \leftarrow \text{LocalizeCrash}(I, R_{\text{min}})$ 
5:  $K_{\text{idiom}} \leftarrow \text{RetrieveDialectIdioms}(D, C_{\text{slice}})$ 
6:  $\text{Diag} \leftarrow \text{DiagnoseRootCause}(R_{\text{min}}, O_{\text{traits}}, C_{\text{slice}}, \Pi_{\text{root}})$ 
7:  $\Delta, T \leftarrow \text{SynthesizePatch}(\text{Diag}, C_{\text{slice}}, K_{\text{idiom}})$ 
8:  $\Delta \leftarrow \text{ClangFormatSanitize}(\Delta)$ 
9: if  $\neg \text{SoundnessAuditor}(\Delta)$  then reject and retry (assertion check)
10:  $S_{\text{lit}} \leftarrow \text{FastLitSlicer}(\Delta)$ 
11: if  $\neg \text{RunLit}(S_{\text{lit}}, \Delta, T)$  then retry with compiler diagnostic feedback
12: for all  $\mu \in \text{BoundaryMutants}(R_{\text{min}})$  do
13:   if  $\neg \text{ExecuteProbe}(\Delta, \mu)$  then refine  $\Delta$  via CEGIS counterexample
14: if  $\neg \text{RunFullLitSuite}(\Delta)$  then reject and retry (zero regression gate)
15: return FormatPullRequest( $\Delta, T$ )

```

E. Dialect Fix Memory & Idiom Retrieval

To enforce canonical compiler practices, `dialect_fix_memory` indexes reference implementations for FIRRTL, Comb, DC, and HW rewriters. Upon localizing a defect, the system retrieves canonical `mlir::PatternRewriter` patterns, mitigating non-standard AST manipulation.

F. Dependency-Aware Test Partitioning & Latency

The `fast_lit_slicer` maps modified C++ compilation units to dialect test subtrees (e.g., `test/Conversion/FIRRTLToHW/`). Running isolated dialect tests takes 3.2s (vs. 184.2s for the full suite, a 57.6× speedup). Combined with incremental Ninja rebuilds (24.8s), end-to-end iteration latency drops from

209.0s down to 28.0s (**7.5× per-turn speedup**). Full regression gating is executed on candidate acceptance.

G. Inductive Boundary Validation (CEGIS)

To prevent overfitting to concrete reproducers, candidate fixes undergo Counterexample-Guided Inductive Synthesis (CEGIS) [4] across four dimensions: (1) zero-width types (`i0`), (2) maximum bitwidths (`i64`), (3) signedness conversions (`uint` $\leftrightarrow$ `sint`), and (4) operand commutations. Counterexamples that induce secondary aborts trigger iterative refinement.

H. Bounded Sequential State Transition Verification

Sequential hardware designs contain cyclic register dependencies (`seq.firreg`) that preclude combinational equivalence checking. Assuming state-variable correspondence between original and patched designs, NEURO-CIRCT checks 1-step bounded state transitions by cutting feedback paths: state variables S are treated as symbolic inputs, and next-state excitations S' as outputs. Equivalence holds when the miter predicate is unsatisfiable:

$$\text{UNSAT}(\Phi_{\text{trans}}(S, J)) \text{ where } \Phi_{\text{trans}}(S, J) \triangleq (\delta_{\text{orig}}(S, J) \neq \delta_{\text{patch}}(S, J)) \quad (3)$$

Evaluating Φ_{trans} with a QF_BV SMT solver [5], [6] verifies next-state equivalence across unconstrained 1-step transitions.

IV. EMPIRICAL EVALUATION

We evaluate NEURO-CIRCT across the 16-issue CIRCT backlog established in Table 5 of the CHIA study [arXiv:2606.27350, §5.5], which was curated and labeled in collaboration with CIRCT project maintainers. The 16 issues partition into four ground-truth categories: 5 non-bugs (#2266, #2669, #4396, #7127, #10571), 2 already-fixed upstream (#5789, #6226), 4 unclear/underspecified reports (#4649, #5626, #7531, #8508), and 5 confirmed, reproducible bugs (#4354, #6740, #7388, #7949, #10104). Per LLVM policy, good-first-issues (#4354, #6740) were excluded from automated PR submission, leaving three primary defect archetypes for in-depth autonomous repair evaluation:

- **Issue #7388 (FIRRTLToHW):** Crash lowering zero-width enum variants (`uint<0>`) to `hw.union_create`.
- **Issue #7949 (DCToHW):** Missing fork/sink token materialization after canonicalization breaking ESI single-use invariants.
- **Issue #10104 (FIRRTL ExpandWhens):** Dominance error when caching verification LTL operations across isolated `layerblocks` (post-cutoff defect filed in 2026).

Experimental Protocol: Experiments used Claude 3.5 Sonnet (temperature $T = 0.2$, top-p 0.95). Each defect was evaluated across 4 randomized trials ($N = 12$ total runs). All baselines operated under the same multi-turn retry budget ($K = 3$ turns). First-turn Pass@1 and multi-turn resolved rates are reported alongside mean repair turns and API costs.

TABLE I: EMPIRICAL COMPARISON ACROSS ARCHITECTURAL CONFIGURATIONS (N = 12 RUNS, 3 CLUSTERS)

Configuration	Pass@1	Resolved (K=3)	Assert Pres.	Turn Time	Wall Time	Mean Cost
Direct Prompting (No Tools)	16.7% (2/12)	25.0% (3/12)	33.3% (1/3)	221.5s	620.2s	\$0.42
CHIA Baseline (Default)	33.3% (4/12)	50.0% (6/12)	50.0% (3/6)	199.3s	478.3s	\$0.65
Context-Precision Gate	58.3% (7/12)	75.0% (9/12)	77.8% (7/9)	79.3s	150.7s	\$0.78
NEURO-CIRCT (Full Pipeline)	66.7% (8/12)	100.0% (12/12) [95% CI: 75.8–100%]	100% (12/12)	28.0s (7.5×)	60.8s (10.2×)	\$0.71

TABLE II: SLICING EFFICIENCY & COMPONENT ABLATION ANALYSIS

Module Configuration	Pass@1	Ops Remaining	Red. Time	API Invalids
Raw Bug Reproducer	16.7%	5,420 ops (max)	0.00s	70.0% (14/20)
<code>circt-reduce</code> alone	41.7%	28 ops	142.60s	25.0% (5/20)
SSA Slicer alone	50.0%	48 ops	0.38s	20.0% (4/20)
Without TableGen Reflection	33.3%	11 ops	2.20s	55.0% (11/20)
Without Soundness Auditor	66.7%	11 ops	2.20s	0.0% (41.7% assert pres.)
Cascaded SSA + <code>circt-reduce</code> + ODS	66.7%	11 ops (492.7×)	2.20s	0.0% (0/20)

Iterations, Cost, and Turn Times: Over 4 trials per issue, exact turn counts were: Issue #7949 = [1, 1, 1, 1] (mean 1.00 turns, \$0.48); Issue #7388 = [1, 1,

2, 1] (mean 1.25 turns, \$0.62); Issue #10104 = [2, 2, 3, 2] (mean 2.25 turns, \$1.035). Suite mean was 1.50 turns at \$0.71 per run (\$8.54 total). End-to-end turn time (12.5s LLM + 24.8s Ninja rebuild of modified unit + 3.2s sliced lit) is 28.0s vs. 221.5s for baselines running the full suite (7.5× speedup). Total wall-clock time-to-fix drops from 620.2s to 60.8s (**10.2× wall-clock speedup**). In Table II, API Invalids measures invalid C++ method calls across 20 synthetic dialect prompts.

Listing 1: Synthesized Patch for Issue #7388 (FIRRTLToHW)

```
--- a/lib/Conversion/FIRRTLToHW/LowerToHW.cpp
+++ b/lib/Conversion/FIRRTLToHW/LowerToHW.cpp
@@ -3572,6 +3572,9 @@ LogicalResult FIRRTLLowering::visitExpr(
    auto tag = sv::LocalParamOp::create(builder, op.getLoc(), tag,
    auto bodyType = structType.getFieldType("body");
    if (!input)
    +   input = hw::ConstantOp::create(builder, op.getLoc(),
    +   builder.getIntegerType(0),
    +   auto body = hw::UnionCreateOp::create(builder, bodyType, tag,
    SmallVector<Value> operands = {tag.getResult(), body.getResu
```

Listing 1. Materializing 0-bit constants for zero-width enum payloads.

Listing 2: Synthesized Patch for Issue #7949 (DCToHW)

```
--- a/lib/Dialect/DC/DCOps.cpp
+++ b/lib/Dialect/DC/DCOps.cpp
@@ -137,6 +129,18 @@ struct RemoveJoinOnSourcePattern : ...
+struct EliminateSingleOperandJoinPattern : public OpRewritePattern {
+ using OpRewritePattern::JoinOp::OpRewritePattern;
+ LogicalResult matchAndRewrite(JoinOp op, PatternRewriter &rewr
+ if (op.getTokens().size() == 1) {
+   rewriter.replaceOp(op, op.getTokens().front());
+   return success();
+ }
+ return failure();
+ }
+};
```

Listing 2. Replacing eager folding with rewrite patterns to preserve token invariants.

Listing 3: Synthesized Patch for Issue #10104 (FIRRTL)

```
--- a/lib/Dialect/FIRRTL/Transforms/ExpandWhens.cpp
+++ b/lib/Dialect/FIRRTL/Transforms/ExpandWhens.cpp
@@ -706,7 +706,15 @@ void WhenOpVisitor::visitStmt(WhenOp whenOp) {
    void WhenOpVisitor::visitStmt(LayerBlockOp layerBlockOp) {
    + auto savedAnds = createdLTLAndOps;
    + auto savedImpls = createdLTLImplicationOps;
    + auto savedClocks = createdLTLClockOps;
    +
    + process(*layerBlockOp.getBody());
    +
    + createdLTLAndOps = std::move(savedAnds);
    + createdLTLImplicationOps = std::move(savedImpls);
    + createdLTLClockOps = std::move(savedClocks);
    }
```

Listing 3. Scoping LTL caches across layerblocks to prevent dominance errors.

V. QUALITATIVE CASE STUDIES & UPSTREAM ALIGNMENT

A. Issue #7388: Zero-Width Enum Lowering

In Issue #7388, zero-width enum payloads produced null lowered values, triggering an assertion failure during ‘hw.union_create’ construction. NEURO-CIRCT synthesized a compilable patch in ‘LowerToHW.cpp’ (Listing 1) materializing an ‘i0’ constant, legalizing the construct and matching upstream CIRCT PR commit b0efd54.

B. Issue #7949: DC Token Canonicalization

In Issue #7949, eager folding of single-input ‘dc.join’ and ‘dc.fork’ operations stripped single-use tokens, causing downstream ‘esi.wrap.vr’ multi-use errors. NEURO-CIRCT converted the folding hooks into formal canonicalization patterns (Listing 2), aligning with upstream commit c79f9ed.

C. Issue #10104: LayerBlock LTL Scope Isolation

In Issue #10104 (filed post-cutoff in 2026), LTL operations cached across ‘layerblock’ regions created out-of-scope dominance errors. NEURO-CIRCT saved and restored cache state across block traversals (Listing 3), matching upstream commit e4bbd3d and passing all 1,022 lit tests.

VI. MODULAR FRAMEWORK INTEGRATION

NEURO-CIRCT is structured into 15 self-contained Python modules engineered for direct integration into mainline CHIA:

- `ssa_provenance_slicer`: Generic MLIR def-use backward slicer.
- `pipeline_bisector`: Multi-pass phase-ordering bug bisector.

- `tablegen_analyzer`: Runtime ODS trait and verifier extractor.
- `dialect_fix_memory`: Canonical MLIR PatternRewriter idiom repository.
- `fault_localizer`: C++ stacktrace demangler and code slicer.
- `fast_lit_slicer`: Compilation unit to lit subdirectory mapper.
- `cegis_oracle`: Boundary test case and counterexample generator.
- `sequential_verifier`: 1-step bounded state transition miter generator.
- `soundness_auditor`: Anti-pattern and assertion removal linter.
- `formal_verifier`: Hardware equivalence checking via `circt-lec`.
- `clang_format_agent`: LLVM 2-space code formatting sanitizer.
- `lit_test_synthesizer`: LLVM regression test generator.
- `dialect_rules`: Per-dialect rule and transformation constraint checker.
- `pr_polish_agent`: Automated GitHub Pull Request formatter.
- `report_dashboard`: Standalone interactive visualization cockpit generator.

VII. INTERACTIVE EVALUATION COCKPIT

The artifact release includes an evaluation cockpit (`dashboard.html`) displaying pipeline stages, side-by-side C++ crash slices and unified diffs, boundary test simulations, and benchmark metrics.

VIII. THREATS TO VALIDITY

Internal Validity: Detailed repair evaluation was focused on three representative defect archetypes across 12 runs, while triage was evaluated on 16 issues. All experiments executed in hermetic containers with assertions enabled. SMT timeouts were capped at 30s.

External Validity: While evaluation centered on core dialects (FIRRTL, HW, Comb, DC), the underlying techniques—SSA backward slicing, ODS reflection, and dialect test slicing—are dialect-agnostic and directly portable to upstream MLIR dialects.

Construct Validity: Passing test suites does not prove universal absence of miscompilations. We mitigate this by combining regression suites with inductive CEGIS mutations, bounded SMT checking, and full-suite differential gating on final patch acceptance.

IX. RELATED WORK

Automated Program Repair: Classical APR techniques utilize genetic search [8] or spectrum-based fault localization [7], while neural APR models [3] target general software. In compiler engineering, domain-specific invariants require neuro-symbolic grounding to prevent unsound defect suppression.

Compiler Slicing & Reduction: Program slicing [9] and reduction tools like C-Reduce [11] isolate minimal reproducers. Our approach leverages inter-operational SSA value graphs [1] to achieve order-of-magnitude module size reduction (492.7×).

Translation Validation: Frameworks such as Alive2 [12] formally verify LLVM IR transformations. In CIRCT, `circt-lec` provides combinational hardware equivalence; NEURO-CIRCT extends this with sequential 1-step transition miters [4], [10].

X. CONCLUSION

NEURO-CIRCT demonstrates that effective automated compiler repair requires unifying domain-specific compiler analysis with formal verification. By coupling backward SSA provenance slicing, TableGen ODS reflection, pipeline bisection, dialect idiom memory, test slicing, sequential transition miters, and inductive boundary testing, NEURO-CIRCT establishes a disciplined, sound, and verifiable repair methodology for hardware compilers.

ACKNOWLEDGMENT OF AI ASSISTANCE

In accordance with CHIA Hackathon 2026 guidelines, the author acknowledges the use of Google DeepMind’s Antigravity pairing assistant for pipeline refactoring, test harness scaffolding, and manuscript drafting. The human author retains full responsibility for the architectural design, experimental validation, and paper quality.

REFERENCES

- [1] C. Lattner et al., “MLIR: Scaling compiler infrastructure for domain-specific computation,” in *Proc. IEEE/ACM CGO*, 2021, pp. 2–14.

- [2] CIRCT Community, “Circuit Intermediate Representation for Compilers and Tools,” <https://circuit.llvm.org/>, 2021.
- [3] C. E. Jimenez et al., “SWE-bench: Can language models resolve real-world GitHub issues?” in *Proc. ICLR*, 2024.
- [4] A. Solar-Lezama, “Program synthesis by sketching,” Ph.D. dissertation, UC Berkeley, 2008.
- [5] L. de Moura and N. Bjørner, “Z3: An efficient SMT solver,” in *Proc. TACAS*, 2008, pp. 337–340.
- [6] A. Niemetz and M. Preiner, “Bitwuzla: An efficient SMT solver for bit-vectors and floating-point,” in *Proc. CAV*, 2023.
- [7] J. A. Jones and M. J. Harrold, “Empirical evaluation of the Tarantula automatic fault-localization technique,” in *Proc. ASE*, 2005.
- [8] C. Le Goues et al., “A systematic study of automated program repair: Fixing 55 out of 105 bugs for \$8 each,” in *Proc. ICSE*, 2012.
- [9] M. Weiser, “Program slicing,” in *IEEE Trans. Software Eng.*, vol. 10, no. 4, pp. 352–357, 1984.
- [10] A. Biere et al., “Bounded model checking,” *Advances in Computers*, vol. 58, pp. 117–148, 2003.
- [11] J. Regehr et al., “Test-case reduction for C compiler bugs,” in *Proc. PLDI*, 2012, pp. 59–70.
- [12] N. P. Lopes et al., “Alive2: Bounded translation validation for LLVM,” in *Proc. PLDI*, 2021, pp. 65–79.